

```
# Healthcare Diabetic Data - Exploratory Data Analysis (EDA)  
# Building a Fraud Detection System for Healthcare Management
```

```
import pandas as pd  
import numpy as np  
import matplotlib.pyplot as plt  
import seaborn as sns  
import warnings  
from scipy import stats  
from sklearn.preprocessing import LabelEncoder  
import plotly.express as px  
import plotly.graph_objects as go  
from plotly.subplots import make_subplots
```

```
# Configuration  
warnings.filterwarnings('ignore')  
plt.style.use('seaborn-v0_8')  
sns.set_palette("husl")
```

```
# Set up plotting parameters  
plt.rcParams['figure.figsize'] = (12, 8)  
plt.rcParams['font.size'] = 10
```

```
print("Healthcare EDA Analysis - Fraud Detection System")  
print("=" * 50)
```

```
#  
=====
```

1. DATA LOADING AND INITIAL EXPLORATION

```
#  
=====
```

```
def load_and_explore_data(file_path='diabetic_data.csv'):
```

```
    """  
    Load the diabetic data and perform initial exploration  
    """
```

```
    print("\n1. LOADING AND EXPLORING DATA")  
    print("-" * 30)
```

```
# Load the dataset  
df = pd.read_csv(file_path)
```

```
print(f"Dataset shape: {df.shape}")  
print(f"Total encounters: {len(df):,}")
```

```
# Display basic info  
print("\nDataset Info:")  
print(df.info())
```

```
print("\nFirst 5 rows:")  
print(df.head())
```

```
print("\nColumn names:")  
print(df.columns.tolist())
```

```

    return df

# Load the data
df = load_and_explore_data()

#
=====
=====
# 2. DATA PREPROCESSING AND CLEANING
#
=====
=====

def preprocess_data(df):
    """
    Clean and preprocess the data for analysis
    """
    print("\n2. DATA PREPROCESSING")
    print("-" * 25)

    # Create a copy for processing
    df_clean = df.copy()

    # Handle missing values and '?' entries
    print("Handling missing values and '?' entries...")

    # Replace '?' with NaN
    df_clean = df_clean.replace('?', np.nan)

    # Check missing values
    missing_data = df_clean.isnull().sum()
    missing_percent = (missing_data / len(df_clean)) * 100

    missing_df = pd.DataFrame({
        'Missing Count': missing_data,
        'Missing Percentage': missing_percent
    }).sort_values('Missing Percentage', ascending=False)

    print("\nMissing Data Summary:")
    print(missing_df[missing_df['Missing Count'] > 0])

    # Convert numerical columns
    numerical_cols = ['time_in_hospital', 'num_lab_procedures', 'num_procedures',
                      'num_medications', 'number_outpatient', 'number_emergency',
                      'number_inpatient', 'number_diagnoses']

    for col in numerical_cols:
        if col in df_clean.columns:
            df_clean[col] = pd.to_numeric(df_clean[col], errors='coerce')

    print(f"Cleaned dataset shape: {df_clean.shape}")

    return df_clean

# Preprocess the data

```

```

df_clean = preprocess_data(df)

#
=====
=====
# 3. DESCRIPTIVE STATISTICAL ANALYSIS FOR NUMERICAL FEATURES
#
=====
=====

def descriptive_statistics(df):
    """
    Perform comprehensive descriptive statistical analysis
    """
    print("\n3. DESCRIPTIVE STATISTICAL ANALYSIS")
    print("-" * 35)

    # Identify numerical columns
    numerical_cols = df.select_dtypes(include=[np.number]).columns.tolist()

    if numerical_cols:
        print("Numerical Features Summary:")
        desc_stats = df[numerical_cols].describe()
        print(desc_stats)

        # Additional statistics
        print("\nAdditional Statistics:")
        additional_stats = pd.DataFrame({
            'Skewness': df[numerical_cols].skew(),
            'Kurtosis': df[numerical_cols].kurtosis(),
            'IQR': df[numerical_cols].quantile(0.75) - df[numerical_cols].quantile(0.25)
        })
        print(additional_stats)

        # Correlation matrix
        plt.figure(figsize=(12, 8))
        correlation_matrix = df[numerical_cols].corr()
        sns.heatmap(correlation_matrix, annot=True, cmap='coolwarm', center=0,
                    square=True, linewidths=0.5)
        plt.title('Correlation Matrix of Numerical Features')
        plt.tight_layout()
        plt.show()

        return desc_stats, additional_stats, correlation_matrix
    else:
        print("No numerical columns found in the dataset")
        return None, None, None

# Perform descriptive statistics
desc_stats, additional_stats, corr_matrix = descriptive_statistics(df_clean)

#
=====
=====
# 4. VISUALIZATION OF CATEGORICAL FEATURES (RACE AND GENDER)

```

```

#
=====

def visualize_categorical_features(df):
    """
    Visualize the distribution of race and gender
    """
    print("\n4. CATEGORICAL FEATURES VISUALIZATION")
    print("-" * 37)

    fig, axes = plt.subplots(2, 2, figsize=(15, 12))

    # Race distribution
    if 'race' in df.columns:
        race_counts = df['race'].value_counts()
        print(f"Race Distribution:")
        print(race_counts)

        # Bar plot
        race_counts.plot(kind='bar', ax=axes[0,0])
        axes[0,0].set_title('Distribution of Race')
        axes[0,0].set_xlabel('Race')
        axes[0,0].set_ylabel('Count')
        axes[0,0].tick_params(axis='x', rotation=45)

        # Pie chart
        axes[0,1].pie(race_counts.values, labels=race_counts.index, autopct='%1.1f%%')
        axes[0,1].set_title('Race Distribution (Pie Chart)')

    # Gender distribution
    if 'gender' in df.columns:
        gender_counts = df['gender'].value_counts()
        print(f"\nGender Distribution:")
        print(gender_counts)

        # Bar plot
        gender_counts.plot(kind='bar', ax=axes[1,0])
        axes[1,0].set_title('Distribution of Gender')
        axes[1,0].set_xlabel('Gender')
        axes[1,0].set_ylabel('Count')

        # Pie chart
        axes[1,1].pie(gender_counts.values, labels=gender_counts.index, autopct='%1.1f%%')
        axes[1,1].set_title('Gender Distribution (Pie Chart)')

    plt.tight_layout()
    plt.show()

    return race_counts if 'race' in df.columns else None, gender_counts if 'gender' in df.columns else None

# Visualize categorical features
race_dist, gender_dist = visualize_categorical_features(df_clean)

```

```

#
=====
=====
# 5. RELATIONSHIP BETWEEN READMISSION STATUS AND AGE
#
=====
=====

def analyze_readmission_age_relationship(df):
    """
    Explore the relationship between readmission status and age
    """
    print("\n5. READMISSION STATUS vs AGE ANALYSIS")
    print("-" * 36)

    if 'readmitted' in df.columns and 'age' in df.columns:
        # Cross-tabulation
        readmit_age_crosstab = pd.crosstab(df['age'], df['readmitted'], margins=True)
        print("Readmission by Age Group:")
        print(readmit_age_crosstab)

        # Percentage breakdown
        readmit_age_pct = pd.crosstab(df['age'], df['readmitted'], normalize='index') * 100
        print("\nReadmission Percentage by Age Group:")
        print(readmit_age_pct.round(2))

        # Visualization
        fig, axes = plt.subplots(2, 2, figsize=(16, 12))

        # Stacked bar chart
        readmit_age_crosstab.iloc[:, :-1].plot(kind='bar', stacked=True, ax=axes[0,0])
        axes[0,0].set_title('Readmission Status by Age Group (Stacked)')
        axes[0,0].set_xlabel('Age Group')
        axes[0,0].set_ylabel('Count')
        axes[0,0].legend(title='Readmitted')
        axes[0,0].tick_params(axis='x', rotation=45)

        # Percentage stacked bar chart
        readmit_age_pct.plot(kind='bar', stacked=True, ax=axes[0,1])
        axes[0,1].set_title('Readmission Percentage by Age Group')
        axes[0,1].set_xlabel('Age Group')
        axes[0,1].set_ylabel('Percentage')
        axes[0,1].legend(title='Readmitted')
        axes[0,1].tick_params(axis='x', rotation=45)

        # Heatmap
        sns.heatmap(readmit_age_pct, annot=True, fmt='.1f', cmap='YlOrRd', ax=axes[1,0])
        axes[1,0].set_title('Readmission Rate Heatmap by Age')

        # Calculate 30-day readmission rates by age
        if '<30' in df['readmitted'].values:
            df['readmitted_30'] = df['readmitted'] == '<30'
            readmit_30_by_age = df.groupby('age')['readmitted_30'].agg(['count', 'sum', 'mean'])
            readmit_30_by_age.columns = ['Total', '30-day_Readmits', '30-day_Rate']
            readmit_30_by_age['30-day_Rate'] *= 100

```

```

readmit_30_by_age['30-day_Rate'].plot(kind='bar', ax=axes[1,1], color='red', alpha=0.7)
axes[1,1].set_title('30-Day Readmission Rate by Age Group')
axes[1,1].set_xlabel('Age Group')
axes[1,1].set_ylabel('30-Day Readmission Rate (%)')
axes[1,1].tick_params(axis='x', rotation=45)

print("\n30-Day Readmission Analysis:")
print(readmit_30_by_age)

plt.tight_layout()
plt.show()

return readmit_age_crosstab, readmit_age_pct
else:
    print("Required columns (readmitted, age) not found")
    return None, None

# Analyze readmission-age relationship
readmit_age_cross, readmit_age_pct = analyze_readmission_age_relationship(df_clean)

#
=====
=====
# 6. MEDICATION ANALYSIS
#
=====
=====

def analyze_medications(df):
    """
    Analyze medication changes and total medications
    """
    print("\n6. MEDICATION ANALYSIS")
    print("-" * 21)

    # Total medications analysis
    if 'num_medications' in df.columns:
        print("Total Medications Statistics:")
        print(df['num_medications'].describe())

    fig, axes = plt.subplots(2, 2, figsize=(15, 12))

    # Distribution of total medications
    df['num_medications'].hist(bins=30, ax=axes[0,0], alpha=0.7, edgecolor='black')
    axes[0,0].set_title('Distribution of Total Medications')
    axes[0,0].set_xlabel('Number of Medications')
    axes[0,0].set_ylabel('Frequency')

    # Box plot
    df.boxplot(column='num_medications', ax=axes[0,1])
    axes[0,1].set_title('Box Plot of Total Medications')
    axes[0,1].set_ylabel('Number of Medications')

    # Medication change analysis
    if 'change' in df.columns:
        change_counts = df['change'].value_counts()

```

```

print(f"\nMedication Change Distribution:")
print(change_counts)

# Medication change visualization
change_counts.plot(kind='bar', ax=axes[1,0])
axes[1,0].set_title('Distribution of Medication Changes')
axes[1,0].set_xlabel('Medication Change')
axes[1,0].set_ylabel('Count')

# Pie chart
axes[1,1].pie(change_counts.values, labels=change_counts.index, autopct='% 1.1f%% ')
axes[1,1].set_title('Medication Change Distribution')

plt.tight_layout()
plt.show()

# Relationship between medication change and readmission
if 'change' in df.columns and 'readmitted' in df.columns:
    med_change_readmit = pd.crosstab(df['change'], df['readmitted'], normalize='index') * 100
    print("\nMedication Change vs Readmission (%):")
    print(med_change_readmit.round(2))

    plt.figure(figsize=(10, 6))
    sns.heatmap(med_change_readmit, annot=True, fmt='.1f', cmap='RdYlBu')
    plt.title('Medication Change vs Readmission Rate')
    plt.show()

# Analyze medications
analyze_medications(df_clean)

#
=====
=====
# 7. DIAGNOSIS CATEGORIES ANALYSIS
#
=====
=====

def analyze_diagnoses(df):
    """
    Examine the distribution of diagnosis categories
    """
    print("\n7. DIAGNOSIS CATEGORIES ANALYSIS")
    print("-" * 31)

    diag_cols = ['diag_1', 'diag_2', 'diag_3']

    for diag_col in diag_cols:
        if diag_col in df.columns:
            print(f"\n{diag_col.upper()} Analysis:")

            # Top 10 diagnoses
            top_diag = df[diag_col].value_counts().head(10)
            print(f"Top 10 {diag_col}:")
            print(top_diag)

```

```

    # Visualization
    plt.figure(figsize=(12, 6))
    top_diag.plot(kind='bar')
    plt.title(f'Top 10 {diag_col} Codes')
    plt.xlabel('Diagnosis Code')
    plt.ylabel('Frequency')
    plt.xticks(rotation=45)
    plt.tight_layout()
    plt.show()

# Diabetes-related diagnosis analysis
if 'diag_1' in df.columns:
    print("\nDiabetes-related Primary Diagnoses:")
    diabetes_diag = df[df['diag_1'].str.contains('250', na=False)][['diag_1']].value_counts()
    print(diabetes_diag.head(10))

# Analyze diagnoses
analyze_diagnoses(df_clean)

#
=====
=====
# 8. ADMISSION ANALYSIS
#
=====
=====

def analyze_admissions(df):
    """
    Explore distribution across admission types, sources, and discharge dispositions
    """
    print("\n8. ADMISSION ANALYSIS")
    print("-" * 19)

    admission_cols = ['admission_type_id', 'admission_source_id', 'discharge_disposition_id']

    fig, axes = plt.subplots(2, 2, figsize=(16, 12))
    axes = axes.flatten()

    for i, col in enumerate(admission_cols):
        if col in df.columns and i < 3:
            col_counts = df[col].value_counts().head(10)
            print(f"\n{col.replace('_', ' ').title()} Distribution:")
            print(col_counts)

            col_counts.plot(kind='bar', ax=axes[i])
            axes[i].set_title(f'Distribution of {col.replace("_", " ").title()}')
            axes[i].set_xlabel(col.replace('_', ' ').title())
            axes[i].set_ylabel('Count')
            axes[i].tick_params(axis='x', rotation=45)

# Length of stay analysis
if 'time_in_hospital' in df.columns:
    stay_counts = df['time_in_hospital'].value_counts().sort_index()
    print(f"\nLength of Stay Distribution:")
    print(stay_counts.head(10))

```



```

stay_counts.head(14).plot(kind='bar', ax=axes[3])
axes[3].set_title('Distribution of Length of Stay')
axes[3].set_xlabel('Days in Hospital')
axes[3].set_ylabel('Count')

plt.tight_layout()
plt.show()

# Analyze admissions
analyze_admissions(df_clean)

#
=====
=====
# 9. OUTLIER DETECTION AND VISUALIZATION
#
=====
=====

def detect_outliers(df):
    """
    Identify and visualize outliers in numerical features
    """
    print("\n9. OUTLIER DETECTION")
    print("-" * 18)

    numerical_cols = df.select_dtypes(include=[np.number]).columns.tolist()

    if not numerical_cols:
        print("No numerical columns found for outlier detection")
        return

    # Z-score method
    print("Outlier Detection using Z-score (threshold = 3):")
    z_scores = np.abs(stats.zscore(df[numerical_cols].fillna(df[numerical_cols].median())))
    outlier_counts_z = (z_scores > 3).sum()
    print(outlier_counts_z)

    # IQR method
    print("\nOutlier Detection using IQR method:")
    outlier_counts_iqr = {}

    for col in numerical_cols:
        Q1 = df[col].quantile(0.25)
        Q3 = df[col].quantile(0.75)
        IQR = Q3 - Q1
        lower_bound = Q1 - 1.5 * IQR
        upper_bound = Q3 + 1.5 * IQR
        outliers = ((df[col] < lower_bound) | (df[col] > upper_bound)).sum()
        outlier_counts_iqr[col] = outliers

    outlier_df = pd.DataFrame({
        'IQR_Outliers': outlier_counts_iqr,
        'Z_Score_Outliers': outlier_counts_z
    })

```

```

print(outlier_df)

# Visualization
n_cols = len(numerical_cols)
n_rows = (n_cols + 2) // 3

fig, axes = plt.subplots(n_rows, 3, figsize=(18, 4*n_rows))
axes = axes.flatten() if n_rows > 1 else [axes] if n_rows == 1 else axes

for i, col in enumerate(numerical_cols[:len(axes)]):
    if i < len(axes):
        df.boxplot(column=col, ax=axes[i])
        axes[i].set_title(f'Box Plot: {col}')
        axes[i].set_ylabel(col)

# Hide empty subplots
for i in range(len(numerical_cols), len(axes)):
    axes[i].set_visible(False)

plt.tight_layout()
plt.show()

return outlier_df

# Detect outliers
outlier_summary = detect_outliers(df_clean)

#
=====
=====
# 10. FRAUD DETECTION INSIGHTS
#
=====
=====

def fraud_detection_analysis(df):
    """
    Analyze patterns that could indicate healthcare fraud
    """
    print("\n10. FRAUD DETECTION ANALYSIS")
    print("-" * 27)

    # Suspicious pattern 1: Unusually high number of procedures
    if 'num_procedures' in df.columns:
        high_procedures = df[df['num_procedures'] > df['num_procedures'].quantile(0.95)]
        print(f"Patients with unusually high procedures (>95th percentile): {len(high_procedures)}")
        print(f"Average procedures in this group: {high_procedures['num_procedures'].mean():.2f}")

    # Suspicious pattern 2: Multiple readmissions
    if 'patient_nbr' in df.columns and 'readmitted' in df.columns:
        patient_encounters = df['patient_nbr'].value_counts()
        multiple_encounters = patient_encounters[patient_encounters > 1]
        print(f"\nPatients with multiple encounters: {len(multiple_encounters)}")
        print(f"Average encounters per patient: {patient_encounters.mean():.2f}")

    # Suspicious pattern 3: Unusual medication combinations

```

```

if 'num_medications' in df.columns:
    high_medications = df[df['num_medications'] > df['num_medications'].quantile(0.99)]
    print(f"\nPatients with extremely high medications (>99th percentile): {len(high_medications)}")

# Suspicious pattern 4: Short stays with high readmission
if 'time_in_hospital' in df.columns and 'readmitted' in df.columns:
    short_stay_readmit = df[(df['time_in_hospital'] == 1) & (df['readmitted'] == '<30')]
    print(f"\nShort stays (1 day) with 30-day readmission: {len(short_stay_readmit)}")

# Risk score calculation
risk_factors = []

if 'num_procedures' in df.columns:
    df['high_procedures'] = (df['num_procedures'] > df['num_procedures'].quantile(0.95)).astype(int)
    risk_factors.append('high_procedures')

if 'num_medications' in df.columns:
    df['high_medications'] = (df['num_medications'] > df['num_medications'].quantile(0.95)).astype(int)
    risk_factors.append('high_medications')

if 'time_in_hospital' in df.columns:
    df['short_stay'] = (df['time_in_hospital'] <= 1).astype(int)
    risk_factors.append('short_stay')

if risk_factors:
    df['fraud_risk_score'] = df[risk_factors].sum(axis=1)

    print(f"\nFraud Risk Score Distribution:")
    risk_distribution = df['fraud_risk_score'].value_counts().sort_index()
    print(risk_distribution)

    plt.figure(figsize=(10, 6))
    risk_distribution.plot(kind='bar')
    plt.title('Distribution of Fraud Risk Scores')
    plt.xlabel('Risk Score')
    plt.ylabel('Count')
    plt.show()

# High-risk cases
high_risk = df[df['fraud_risk_score'] >= 2]
print(f"\nHigh-risk cases (score >= 2): {len(high_risk)} ({len(high_risk)/len(df)*100:.2f}%)")

# Fraud detection analysis
fraud_detection_analysis(df_clean)

#
=====
=====
# 11. ANALYSIS REPORT GENERATION
#
=====
=====

def generate_analysis_report(df, desc_stats, corr_matrix, outlier_summary):
    """
    Generate comprehensive EDA report for fraud detection system

```

```

"""
print("\n" + "="*70)
print("HEALTHCARE EDA ANALYSIS REPORT")
print("FRAUD DETECTION SYSTEM CONTEXT")
print("="*70)

print("\n1. EXECUTIVE SUMMARY")
print("-" * 20)
print(f"• Dataset contains {len(df):,} healthcare encounters")
print(f"• Analysis covers {df.shape[1]} features including demographics, medical procedures, and outcomes")
print(f"• Key focus areas: readmission patterns, medication management, and fraud detection indicators")

print("\n2. DATA QUALITY ASSESSMENT")
print("-" * 28)
missing_data = df.isnull().sum().sum()
print(f"• Total missing values: {missing_data:,}")
print(f"• Data completeness: {(df.size - missing_data) / df.size * 100:.1f}%")

print("\n3. KEY FINDINGS")
print("-" * 15)

# Demographics
if 'age' in df.columns:
    elderly_patients = len(df[df['age'].isin(['70-80', '80-90', '90-100'])])
    print(f"• Elderly patients (70+): {elderly_patients:,} ({elderly_patients/len(df)*100:.1f}%")

# Readmissions
if 'readmitted' in df.columns:
    readmit_30 = len(df[df['readmitted'] == '<30'])
    print(f"• 30-day readmissions: {readmit_30:,} ({readmit_30/len(df)*100:.1f}%")

# High-risk indicators
if 'num_medications' in df.columns:
    high_med_patients = len(df[df['num_medications'] > df['num_medications'].quantile(0.95)])
    print(f"• High medication patients (>95th percentile): {high_med_patients:,}")

print("\n4. FRAUD DETECTION IMPLICATIONS")
print("-" * 32)
print("• Patterns identified for potential fraud detection:")
print("  - Unusually high procedure counts")
print("  - Excessive medication prescriptions")
print("  - Frequent readmissions with short stays")
print("  - Outlier patterns in treatment intensity")

print("\n5. RECOMMENDATIONS")
print("-" * 18)
print("• Implement real-time monitoring for high-risk score patients")
print("• Develop predictive models for 30-day readmission prevention")
print("• Create alerts for unusual medication/procedure combinations")
print("• Establish care transition protocols for elderly patients")
print("• Monitor patients with multiple encounters for potential fraud")

print("\n6. TECHNICAL NOTES")
print("-" * 16)
print(f"• Analysis performed using Python pandas, matplotlib, seaborn")
print(f"• Statistical methods: descriptive statistics, correlation analysis, outlier detection")

```

```

print(f"• Fraud risk scoring based on multiple behavioral indicators")

print("\n" + "="*70)
print("END OF ANALYSIS REPORT")
print("="*70)

# Generate final report
generate_analysis_report(df_clean, desc_stats, corr_matrix, outlier_summary)

print("\n🏥 Healthcare EDA Analysis Complete!")
print("All visualizations and analyses have been generated.")
print("The notebook is ready for your fraud detection system implementation.")
Healthcare EDA Analysis - Fraud Detection System
=====

```

1. LOADING AND EXPLORING DATA

Dataset shape: (101766, 50)
Total encounters: 101,766

Dataset Info:

<class 'pandas.core.frame.DataFrame'>

RangeIndex: 101766 entries, 0 to 101765

Data columns (total 50 columns):

#	Column	Non-Null Count	Dtype
0	encounter_id	101766 non-null	int64
1	patient_nbr	101766 non-null	int64
2	race	101766 non-null	object
3	gender	101766 non-null	object
4	age	101766 non-null	object
5	weight	101766 non-null	object
6	admission_type_id	101766 non-null	int64
7	discharge_disposition_id	101766 non-null	int64
8	admission_source_id	101766 non-null	int64
9	time_in_hospital	101766 non-null	int64
10	payer_code	101766 non-null	object
11	medical_specialty	101766 non-null	object
12	num_lab_procedures	101766 non-null	int64
13	num_procedures	101766 non-null	int64
14	num_medications	101766 non-null	int64
15	number_outpatient	101766 non-null	int64
16	number_emergency	101766 non-null	int64
17	number_inpatient	101766 non-null	int64
18	diag_1	101766 non-null	object
19	diag_2	101766 non-null	object
20	diag_3	101766 non-null	object
21	number_diagnoses	101766 non-null	int64
22	max_glu_serum	5346 non-null	object
23	A1Cresult	17018 non-null	object
24	metformin	101766 non-null	object
25	repaglinide	101766 non-null	object
26	nateglinide	101766 non-null	object
27	chlorpropamide	101766 non-null	object
28	glimepiride	101766 non-null	object
29	acetoexamide	101766 non-null	object

30	glipizide	101766 non-null object
31	glyburide	101766 non-null object
32	tolbutamide	101766 non-null object
33	pioglitazone	101766 non-null object
34	rosiglitazone	101766 non-null object
35	acarbose	101766 non-null object
36	miglitol	101766 non-null object
37	troglitazone	101766 non-null object
38	tolazamide	101766 non-null object
39	examide	101766 non-null object
40	citoglipton	101766 non-null object
41	insulin	101766 non-null object
42	glyburide-metformin	101766 non-null object
43	glipizide-metformin	101766 non-null object
44	glimepiride-pioglitazone	101766 non-null object
45	metformin-rosiglitazone	101766 non-null object
46	metformin-pioglitazone	101766 non-null object
47	change	101766 non-null object
48	diabetesMed	101766 non-null object
49	readmitted	101766 non-null object

dtypes: int64(13), object(37)

memory usage: 38.8+ MB

None

First 5 rows:

	encounter_id	patient_nbr	race	gender	age	weight	\
0	2278392	8222157	Caucasian	Female	[0-10)	?	
1	149190	55629189	Caucasian	Female	[10-20)	?	
2	64410	86047875	AfricanAmerican	Female	[20-30)	?	
3	500364	82442376	Caucasian	Male	[30-40)	?	
4	16680	42519267	Caucasian	Male	[40-50)	?	

	admission_type_id	discharge_disposition_id	admission_source_id	\
0	6	25	1	
1	1	1	7	
2	1	1	7	
3	1	1	7	
4	1	1	7	

	time_in_hospital	...	citoglipton	insulin	glyburide-metformin	\
0	1	...	No	No	No	
1	3	...	No	Up	No	
2	2	...	No	No	No	
3	2	...	No	Up	No	
4	1	...	No	Steady	No	

	glipizide-metformin	glimepiride-pioglitazone	metformin-rosiglitazone	\
0	No	No	No	
1	No	No	No	
2	No	No	No	
3	No	No	No	
4	No	No	No	

	metformin-pioglitazone	change	diabetesMed	readmitted
0	No	No	No	NO
1	No	Ch	Yes	>30

2	No	No	Yes	NO
3	No	Ch	Yes	NO
4	No	Ch	Yes	NO

[5 rows x 50 columns]

Column names:

['encounter_id', 'patient_nbr', 'race', 'gender', 'age', 'weight', 'admission_type_id', 'discharge_disposition_id', 'admission_source_id', 'time_in_hospital', 'payer_code', 'medical_specialty', 'num_lab_procedures', 'num_procedures', 'num_medications', 'number_outpatient', 'number_emergency', 'number_inpatient', 'diag_1', 'diag_2', 'diag_3', 'number_diagnoses', 'max_glu_serum', 'A1Cresult', 'metformin', 'repaglinide', 'nateglinide', 'chlorpropamide', 'glimepiride', 'acetohexamide', 'glipizide', 'glyburide', 'tolbutamide', 'pioglitazone', 'rosiglitazone', 'acarbose', 'miglitol', 'troglitazone', 'tolazamide', 'examide', 'citoglipton', 'insulin', 'glyburide-metformin', 'glipizide-metformin', 'glimepiride-pioglitazone', 'metformin-rosiglitazone', 'metformin-pioglitazone', 'change', 'diabetesMed', 'readmitted']

2. DATA PREPROCESSING

Handling missing values and '?' entries...

Missing Data Summary:

	Missing Count	Missing Percentage
weight	98569	96.858479
max_glu_serum	96420	94.746772
A1Cresult	84748	83.277322
medical_specialty	49949	49.082208
payer_code	40256	39.557416
race	2273	2.233555
diag_3	1423	1.398306
diag_2	358	0.351787
diag_1	21	0.020636

Cleaned dataset shape: (101766, 50)

3. DESCRIPTIVE STATISTICAL ANALYSIS

Numerical Features Summary:

	encounter_id	patient_nbr	admission_type_id \
count	1.017660e+05	1.017660e+05	101766.000000
mean	1.652016e+08	5.433040e+07	2.024006
std	1.026403e+08	3.869636e+07	1.445403
min	1.252200e+04	1.350000e+02	1.000000
25%	8.496119e+07	2.341322e+07	1.000000
50%	1.523890e+08	4.550514e+07	1.000000
75%	2.302709e+08	8.754595e+07	3.000000
max	4.438672e+08	1.895026e+08	8.000000

	discharge_disposition_id	admission_source_id	time_in_hospital \
count	101766.000000	101766.000000	101766.000000
mean	3.715642	5.754437	4.395987
std	5.280166	4.064081	2.985108
min	1.000000	1.000000	1.000000
25%	1.000000	1.000000	2.000000
50%	1.000000	7.000000	4.000000
75%	4.000000	7.000000	6.000000
max	28.000000	25.000000	14.000000

	num_lab_procedures	num_procedures	num_medications	number_outpatient \
--	--------------------	----------------	-----------------	---------------------

count	101766.000000	101766.000000	101766.000000	101766.000000
mean	43.095641	1.339730	16.021844	0.369357
std	19.674362	1.705807	8.127566	1.267265
min	1.000000	0.000000	1.000000	0.000000
25%	31.000000	0.000000	10.000000	0.000000
50%	44.000000	1.000000	15.000000	0.000000
75%	57.000000	2.000000	20.000000	0.000000
max	132.000000	6.000000	81.000000	42.000000

	number_emergency	number_inpatient	number_diagnoses
count	101766.000000	101766.000000	101766.000000
mean	0.197836	0.635566	7.422607
std	0.930472	1.262863	1.933600
min	0.000000	0.000000	1.000000
25%	0.000000	0.000000	6.000000
50%	0.000000	0.000000	8.000000
75%	0.000000	1.000000	9.000000
max	76.000000	21.000000	16.000000

Additional Statistics:

	Skewness	Kurtosis	IQR
encounter_id	0.699142	-0.102071	1.453097e+08
patient_nbr	0.471281	-0.347372	6.413273e+07
admission_type_id	1.591984	1.942476	2.000000e+00
discharge_disposition_id	2.563067	6.003347	3.000000e+00
admission_source_id	1.029935	1.744989	6.000000e+00
time_in_hospital	1.133999	0.850251	4.000000e+00
num_lab_procedures	-0.236544	-0.245074	2.600000e+01
num_procedures	1.316415	0.857110	2.000000e+00
num_medications	1.326672	3.468155	1.000000e+01
number_outpatient	8.832959	147.907736	0.000000e+00
number_emergency	22.855582	1191.686726	0.000000e+00
number_inpatient	3.614139	20.719397	1.000000e+00
number_diagnoses	-0.876746	-0.079056	3.000000e+00

4. CATEGORICAL FEATURES VISUALIZATION

Race Distribution:

```

race
Caucasian      76099
AfricanAmerican 19210
Hispanic        2037
Other           1506
Asian           641
Name: count, dtype: int64

```

Gender Distribution:

```

gender
Female      54708
Male        47055
Unknown/Invalid 3
Name: count, dtype: int64

```

5. READMISSION STATUS vs AGE ANALYSIS

Readmission by Age Group:

readmitted	<30	>30	NO	All
age				
[0-10)	3	26	132	161
[10-20)	40	224	427	691
[20-30)	236	510	911	1657
[30-40)	424	1187	2164	3775
[40-50)	1027	3278	5380	9685
[50-60)	1668	5917	9671	17256
[60-70)	2502	7897	12084	22483
[70-80)	3069	9475	13524	26068
[80-90)	2078	6223	8896	17197
[90-100)	310	808	1675	2793
All	11357	35545	54864	101766

Readmission Percentage by Age Group:

readmitted	<30	>30	NO
age			
[0-10)	1.86	16.15	81.99
[10-20)	5.79	32.42	61.79
[20-30)	14.24	30.78	54.98
[30-40)	11.23	31.44	57.32
[40-50)	10.60	33.85	55.55
[50-60)	9.67	34.29	56.04
[60-70)	11.13	35.12	53.75
[70-80)	11.77	36.35	51.88
[80-90)	12.08	36.19	51.73
[90-100)	11.10	28.93	59.97

30-Day Readmission Analysis:

	Total	30-day_Readmits	30-day_Rate
age			
[0-10)	161	3	1.863354
[10-20)	691	40	5.788712
[20-30)	1657	236	14.242607
[30-40)	3775	424	11.231788
[40-50)	9685	1027	10.604027
[50-60)	17256	1668	9.666203
[60-70)	22483	2502	11.128408
[70-80)	26068	3069	11.773055
[80-90)	17197	2078	12.083503
[90-100)	2793	310	11.099177

6. MEDICATION ANALYSIS

Total Medications Statistics:

count	101766.000000
mean	16.021844
std	8.127566
min	1.000000
25%	10.000000
50%	15.000000
75%	20.000000
max	81.000000

Name: num_medications, dtype: float64

Medication Change Distribution:

change

No 54755

Ch 47011

Name: count, dtype: int64

Medication Change vs Readmission (%):

readmitted <30 >30 NO

change

Ch 11.82 36.74 51.44

No 10.59 33.37 56.04

7. DIAGNOSIS CATEGORIES ANALYSIS

DIAG_1 Analysis:

Top 10 diag_1:

diag_1

428 6862

414 6581

786 4016

410 3614

486 3508

427 2766

491 2275

715 2151

682 2042

434 2028

Name: count, dtype: int64

DIAG_2 Analysis:

Top 10 diag_2:

diag_2

276 6752

428 6662

250 6071

427 5036

401 3736

496 3305

599 3288

403 2823

414 2650

411 2566

Name: count, dtype: int64

DIAG_3 Analysis:

Top 10 diag_3:

diag_3

250 11555

401 8289

276 5175

428 4577

427 3955

414 3664

496 2605

```
403  2357
585  1992
272  1969
Name: count, dtype: int64
```

Diabetes-related Primary Diagnoses:

```
diag_1
250.8  1680
250.6  1183
250.7   871
250.13  851
250.02  675
250.11  625
250.12  417
250.82  412
250.1   313
250.4   267
Name: count, dtype: int64
```

8. ADMISSION ANALYSIS

Admission Type Id Distribution:

```
admission_type_id
1  53990
3  18869
2  18480
6  5291
5  4785
8   320
7    21
4     10
Name: count, dtype: int64
```

Admission Source Id Distribution:

```
admission_source_id
7  57494
1  29565
17  6781
4  3187
6  2264
2  1104
5   855
3   187
20  161
9   125
Name: count, dtype: int64
```

Discharge Disposition Id Distribution:

```
discharge_disposition_id
1  60234
3  13954
6  12902
18  3691
2   2128
22  1993
```

```

11  1642
5   1184
25  989
4   815
Name: count, dtype: int64

```

Length of Stay Distribution:

```
time_in_hospital
```

```

1   14208
2   17224
3   17756
4   13924
5    9966
6    7539
7    5859
8    4391
9    3002
10   2342

```

```
Name: count, dtype: int64
```

9. OUTLIER DETECTION

Outlier Detection using Z-score (threshold = 3):

```

encounter_id      0
patient_nbr       847
admission_type_id 341
discharge_disposition_id 3588
admission_source_id 175
time_in_hospital  1042
num_lab_procedures 43
num_procedures     0
num_medications    1361
number_outpatient   1457
number_emergency    1664
number_inpatient    2016
number_diagnoses    281
dtype: int64

```

Outlier Detection using IQR method:

	IQR_Outliers	Z_Score_Outliers
encounter_id	0	0
patient_nbr	247	847
admission_type_id	341	341
discharge_disposition_id	9818	3588
admission_source_id	6956	175
time_in_hospital	2252	1042
num_lab_procedures	143	43
num_procedures	4954	0
num_medications	2557	1361
number_outpatient	16739	1457
number_emergency	11383	1664
number_inpatient	7049	2016
number_diagnoses	281	281

10. FRAUD DETECTION ANALYSIS

Patients with unusually high procedures (>95th percentile): 4954
Average procedures in this group: 6.00

Patients with multiple encounters: 16773
Average encounters per patient: 1.42

Patients with extremely high medications (>99th percentile): 960

Short stays (1 day) with 30-day readmission: 1162

Fraud Risk Score Distribution:

fraud_risk_score

0 80266

1 19316

2 2181

3 3

Name: count, dtype: int64

High-risk cases (score >= 2): 2184 (2.15%)

=====

HEALTHCARE EDA ANALYSIS REPORT

FRAUD DETECTION SYSTEM CONTEXT

=====

1. EXECUTIVE SUMMARY

- Dataset contains 101,766 healthcare encounters
- Analysis covers 55 features including demographics, medical procedures, and outcomes
- Key focus areas: readmission patterns, medication management, and fraud detection indicators

2. DATA QUALITY ASSESSMENT

- Total missing values: 374,017
- Data completeness: 93.3%

3. KEY FINDINGS

- Elderly patients (70+): 46,058 (45.3%)
- 30-day readmissions: 11,357 (11.2%)
- High medication patients (>95th percentile): 4,525

4. FRAUD DETECTION IMPLICATIONS

- Patterns identified for potential fraud detection:
 - Unusually high procedure counts
 - Excessive medication prescriptions
 - Frequent readmissions with short stays
 - Outlier patterns in treatment intensity

5. RECOMMENDATIONS

- Implement real-time monitoring for high-risk score patients
- Develop predictive models for 30-day readmission prevention
- Create alerts for unusual medication/procedure combinations

- Establish care transition protocols for elderly patients
- Monitor patients with multiple encounters for potential fraud

6. TECHNICAL NOTES

- Analysis performed using Python pandas, matplotlib, seaborn
- Statistical methods: descriptive statistics, correlation analysis, outlier detection
- Fraud risk scoring based on multiple behavioral indicators

=====

END OF ANALYSIS REPORT

=====

🎉 Healthcare EDA Analysis Complete!

All visualizations and analyses have been generated.

The notebook is ready for your fraud detection system implementation.